

Mortality Prediction for Acute Respiratory Failure Using MIMIC-IV

```
In [ ]: """
Import Libraries
"""

import pandas as pd
import numpy as np
import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, roc_auc_score, precision_score, recall_
from imblearn.over_sampling import SMOTE
from sklearn.utils.class_weight import compute_class_weight
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from xgboost import XGBClassifier
import warnings

# Suppress warnings
warnings.filterwarnings('ignore')
```

```
In [ ]: # Mount Google Drive
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

MIMIC-IV: Loading and Filter Data for Acute Respiratory Failure Analysis

```
In [ ]: """
Common function to read mimic csv
"""

def read_mimic_csv_file(mimic_csv_file_name: str, low_memory: bool = False, chunksize: int = 10000) -> pd.DataFrame:
    """
    Read a CSV file from the MIMIC-IV dataset into a pandas DataFrame.

    Parameters:
    - mimic_csv_file_name (str): Name of the CSV file.
    - low_memory (bool): Whether to use low memory mode when reading.
    - chunksize (int, optional): Number of rows per chunk if reading in chunks.
    """
```

```

Returns:
- pd.DataFrame
"""

# Define the root directory of MIMIC-IV data in Google Drive
mimic_root_dir_path = "/content/drive/MyDrive/Colab Notebooks/AIH/MIMIC-IV/"
file_path = mimic_root_dir_path + mimic_csv_file_name

return pd.read_csv(file_path, low_memory=low_memory, chunksize=chunksize)

```

```

In [ ]: """
    Read, filter, and merge MIMIC-IV data for acute respiratory failure patients.
    """

# Load diagnoses data
arf_diagnoses_df = read_mimic_csv_file("diagnoses_icd.csv.gz")

# Define relevant ICD-9 and ICD-10 codes for acute respiratory failure(MIMIC-IV con
arf_icd_codes = {'51851', '51881', 'J960', 'J9600', 'J9601', 'J9602'}

# Filter diagnoses dataset
arf_diagnoses_df = arf_diagnoses_df[arf_diagnoses_df['icd_code'].isin(arf_icd_codes

# Drop unnecessary columns
arf_diagnoses_df.drop(columns=['seq_num', 'icd_code', 'icd_version'], inplace=True,

# Remove duplicates
arf_diagnoses_df.drop_duplicates(inplace=True)

# Merge with admissions data
arf_admissions_df = read_mimic_csv_file('admissions.csv.gz')

arf_merged_df = arf_diagnoses_df.merge(
    arf_admissions_df, on=['subject_id', 'hadm_id'], how='inner'
)

arf_merged_df.drop(columns=['dischtime', 'deathtime', 'admit_provider_id', 'dischar
    'language', 'edregtime', 'edouttime'], inplace=True, er

arf_merged_df.drop_duplicates(inplace=True)
arf_merged_df.reset_index(drop=True, inplace=True)

# Merge with patient demographics
arf_patients_df = read_mimic_csv_file('patients.csv.gz')

arf_merged_df = arf_merged_df.merge(
    arf_patients_df, on=['subject_id'], how='inner'
)

arf_merged_df.drop(columns=['dod', 'anchor_year_group'], inplace=True, errors='igno
arf_merged_df.drop_duplicates(inplace=True)
arf_merged_df.reset_index(drop=True, inplace=True)

# Merge with ICU stays
arf_icustays_df = read_mimic_csv_file('icustays.csv.gz')

```

```

arf_merged_df = arf_merged_df.merge(
    arf_icustays_df, on=['subject_id', 'hadm_id'], how='inner'
)

arf_merged_df.drop(columns=['last_careunit', 'intime', 'outtime', 'los', 'stay_id'])
arf_merged_df.drop_duplicates(inplace=True)
arf_merged_df.reset_index(drop=True, inplace=True)

# Define lab test keywords related to respiratory function
resp_lab_tests = {
    'oxygen saturation', 'oxygen', 'ph', 'pco2',
    'bicarbonate', 'lactate', 'calculated bicarbonate, whole blood'
}

# Load lab item details
lab_items_df = read_mimic_csv_file('d_labitems.csv.gz')

# Filter respiratory-related blood lab items
lab_items_df = lab_items_df[
    (lab_items_df['fluid'] == 'Blood') &
    (lab_items_df['label'].str.lower().str.strip().isin(resp_lab_tests))
].copy()

# Drop unnecessary columns
lab_items_df.drop(columns=['fluid', 'category'], inplace=True, errors='ignore')
lab_items_df.drop_duplicates(inplace=True)
lab_items_df.reset_index(drop=True, inplace=True)

# Extract unique subject_id and hadm_id pairs
subject_hadm_set = arf_merged_df[['subject_id', 'hadm_id']].drop_duplicates().reset_index()

# Process lab events data in chunks to manage memory efficiently
lab_chunks = []
for lab_chunk in read_mimic_csv_file('labevents.csv.gz', low_memory=False, chunksize=10000):
    # Drop irrelevant columns
    lab_chunk.drop(columns=['labevent_id', 'value', 'valueuom', 'flag', 'ref_range',
                           'priority', 'specimen_id', 'order_provider_id', 'storet'],
                  inplace=True, errors='ignore')

    # Merge with filtered lab items
    lab_chunk = lab_chunk.merge(lab_items_df, on='itemid', how='inner')
    lab_chunk.drop(columns=['itemid'], inplace=True, errors='ignore')

    # Keep only data for acute respiratory failure patients
    lab_chunk = lab_chunk.merge(subject_hadm_set, on=['subject_id', 'hadm_id'], how='inner')

    # Sort for time-based aggregation
    lab_chunk.sort_values(by=['subject_id', 'hadm_id', 'charttime'], inplace=True)

    # Aggregate lab test values by median per subject_id, hadm_id, and label
    lab_chunk = lab_chunk.groupby(['subject_id', 'hadm_id', 'label'], as_index=False).median()

    lab_chunks.append(lab_chunk)

# Merge processed lab event data with the main dataset
if lab_chunks:

```

```
arf_merged_df = arf_merged_df.merge(pd.concat(lab_chunks, ignore_index=True),
                                     on=['subject_id', 'hadm_id'], how='inner')

# Remove duplicate rows
arf_merged_df.drop_duplicates(subset=['subject_id', 'hadm_id', 'label'], inplace=True)
arf_merged_df.reset_index(drop=True, inplace=True)
```

```
In [ ]: # Display dataset info
arf_merged_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 83237 entries, 0 to 83236
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   subject_id            83237 non-null  int64
1   hadm_id               83237 non-null  int64
2   admittime             83237 non-null  object
3   admission_type        83237 non-null  object
4   admission_location    83237 non-null  object
5   insurance             81768 non-null  object
6   marital_status        72965 non-null  object
7   race                  83237 non-null  object
8   hospital_expire_flag  83237 non-null  int64
9   gender                83237 non-null  object
10  anchor_age            83237 non-null  int64
11  anchor_year           83237 non-null  int64
12  first_careunit        83237 non-null  object
13  label                 83237 non-null  object
14  valuenum              83213 non-null  float64
dtypes: float64(1), int64(5), object(9)
memory usage: 9.5+ MB
```

```
In [ ]: arf_merged_df.head()
```

Out[]:

	subject_id	hadm_id	admittime	admission_type	admission_location	insurance	marital
0	10000690	25860671	2150-11-02 18:02:00	EW EMER.	EMERGENCY ROOM	Medicare	WIE
1	10000690	25860671	2150-11-02 18:02:00	EW EMER.	EMERGENCY ROOM	Medicare	WIE
2	10000690	25860671	2150-11-02 18:02:00	EW EMER.	EMERGENCY ROOM	Medicare	WIE
3	10001843	26133978	2134-12-05 00:10:00	URGENT	TRANSFER FROM HOSPITAL	Medicare	
4	10001843	26133978	2134-12-05 00:10:00	URGENT	TRANSFER FROM HOSPITAL	Medicare	

Preprocessing and Feature Engineering for Acute Respiratory Failure Analysis

In []:

```
# Check Missing Values
arf_merged_df.isnull().sum()
```

```
Out[ ]:
```

subject_id	0
hadm_id	0
admittime	0
admission_type	0
admission_location	0
insurance	1469
marital_status	10272
race	0
hospital_expire_flag	0
gender	0
anchor_age	0
anchor_year	0
first_careunit	0
label	0
valuenum	24

dtype: int64

```
In [ ]: # Get unique values in the 'marital_status' column
arf_merged_df['marital_status'].unique()
```

```
Out[ ]: array(['WIDOWED', 'SINGLE', 'MARRIED', nan, 'DIVORCED'], dtype=object)
```

```
In [ ]: # Get unique values in the 'first_careunit' column
arf_merged_df['first_careunit'].unique()
```

```
Out[ ]: array(['Medical Intensive Care Unit (MICU)',
               'Medical/Surgical Intensive Care Unit (MICU/SICU)',
               'Coronary Care Unit (CCU)', 'Trauma SICU (TSICU)',
               'Neuro Surgical Intensive Care Unit (Neuro SICU)',
               'Surgical Intensive Care Unit (SICU)', 'Neuro Intermediate',
               'Cardiac Vascular Intensive Care Unit (CVICU)', 'Neuro Stepdown',
               'Intensive Care Unit (ICU)', 'Surgery/Vascular/Intermediate',
               'PACU', 'Medicine', 'Surgery/Trauma', 'Med/Surg'], dtype=object)
```

```
In [ ]: # Get unique values in the 'race' column
arf_merged_df['race'].unique()
```

```
Out[ ]: array(['WHITE', 'BLACK/AFRICAN AMERICAN', 'UNKNOWN', 'PORTUGUESE',
              'BLACK/CAPE VERDEAN', 'ASIAN - SOUTH EAST ASIAN',
              'WHITE - OTHER EUROPEAN', 'WHITE - BRAZILIAN', 'UNABLE TO OBTAIN',
              'HISPANIC/LATINO - CUBAN', 'HISPANIC OR LATINO',
              'HISPANIC/LATINO - DOMINICAN', 'HISPANIC/LATINO - PUERTO RICAN',
              'ASIAN - CHINESE', 'OTHER',
              'NATIVE HAWAIIAN OR OTHER PACIFIC ISLANDER',
              'HISPANIC/LATINO - SALVADORAN', 'BLACK/CARIBBEAN ISLAND', 'ASIAN',
              'ASIAN - ASIAN INDIAN', 'HISPANIC/LATINO - HONDURAN',
              'HISPANIC/LATINO - COLUMBIAN', 'WHITE - RUSSIAN',
              'PATIENT DECLINED TO ANSWER', 'BLACK/AFRICAN',
              'HISPANIC/LATINO - CENTRAL AMERICAN', 'ASIAN - KOREAN',
              'SOUTH AMERICAN', 'WHITE - EASTERN EUROPEAN',
              'AMERICAN INDIAN/ALASKA NATIVE', 'HISPANIC/LATINO - GUATEMALAN',
              'HISPANIC/LATINO - MEXICAN', 'MULTIPLE RACE/ETHNICITY'],
              dtype=object)
```

```
In [ ]: # Get statistical summary of 'anchor_age' to understand its distribution for approp
arf_merged_df['anchor_age'].describe()
```

```
Out[ ]:
```

	anchor_age
count	83237.000000
mean	64.066869
std	16.117273
min	18.000000
25%	54.000000
50%	65.000000
75%	76.000000
max	91.000000

dtype: float64

```
In [ ]: import itertools

"""
    Preprocessing and Feature Engineering for required features
"""

# Create a copy of the merged Acute Respiratory Failure dataset for processing
arf_processed_df = arf_merged_df.copy()

# Map Gender Column
arf_processed_df['gender'] = arf_processed_df['gender'].map({'F': 'Female', 'M': 'M

# Handle missing values in marital status by replacing NaNs with 'Unknown'
arf_processed_df['marital_status'] = arf_processed_df['marital_status'].fillna('Unk

# Handle missing values in insurance by replacing NaNs with 'Unknown'
```

```

arf_processed_df['insurance'] = arf_processed_df['insurance'].fillna('Unknown')

# Convert admission time to datetime format
arf_processed_df['admittime'] = pd.to_datetime(arf_processed_df['admittime'])

# Compute patient age at admission using MIMIC-IV anchor values
arf_processed_df['admission_age'] = (
    arf_processed_df['anchor_age'] +
    (arf_processed_df['admittime'].dt.year - arf_processed_df['anchor_year'])
)

# Categorize patients into age groups: Young (<30), Adult (30-60), Senior (60+)
arf_processed_df['age_group'] = pd.cut(
    arf_processed_df['admission_age'],
    bins=[0, 30, 60, float('inf')],
    labels=['Young', 'Adult', 'Senior'],
    right=False
)

# Remove unnecessary columns after computing age group
arf_processed_df.drop(columns=['admittime', 'anchor_year', 'anchor_age', 'admission

# Convert age group to string type
arf_processed_df['age_group'] = arf_processed_df['age_group'].astype(str)

# Standardize race categories by grouping similar values
arf_processed_df['race'] = arf_processed_df['race'].replace(
    {r"ASIAN\D*": "ASIAN",
     r"WHITE\D*": "WHITE",
     r"HISPANIC\D*": "HISPANIC/LATINO",
     r"BLACK\D*": "BLACK/AFRICAN AMERICAN"},
    regex=True
)

# Replace ambiguous race values with 'OTHER/UNKNOWN'
arf_processed_df['race'] = arf_processed_df['race'].replace(
    ['UNABLE TO OBTAIN', 'OTHER', 'PATIENT DECLINED TO ANSWER', 'UNKNOWN', 'MULTIPL
    'OTHER/UNKNOWN'
)

# Standardize ICU (first care unit) categories by grouping related units
arf_processed_df['first_careunit'] = arf_processed_df['first_careunit'].replace(
    {r"Medical/Surgical\D*": "MICU, SICU",
     r"Medical\D*": "MICU",
     r"Neuro\D*": "NSICU",
     r"Cardiac\D*": "CVICU",
     r"Coronary\D*": "CCU",
     r"Trauma SICU\D*": "TSICU",
     r"Surgical\D*": "SICU",
     r"Intensive Care Unit\D*": "ICU"},
    regex=True
)

# Convert uncommon ICU categories into 'OTHERICU'
arf_processed_df['first_careunit'] = arf_processed_df['first_careunit'].replace(
    ['Surgery/Vascular/Intermediate', 'PACU', 'Medicine', 'Surgery/Trauma', 'Med/Su

```


Out[]:

	hospital_expire_flag	CCU	CVICU	ICU	MICU	NSICU	OTHER_ICU	SICU	TSICU	Bicarb
0	0	0	0	0	1	0	0	0	0	0
1	1	0	0	0	1	0	0	1	0	0
2	1	0	0	0	1	0	0	0	0	0
3	0	1	0	0	0	0	0	0	0	0
4	0	0	0	0	1	0	0	1	0	0

5 rows × 60 columns



```
In [ ]: # Exploring the columns of preprocessed data
processed_data = arf_processed_df.copy()
processed_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 17003 entries, 0 to 17002
```

```
Data columns (total 60 columns):
```

#	Column	Non-Null Count	Dtype
0	hospital_expire_flag	17003 non-null	int64
1	CCU	17003 non-null	int64
2	CVICU	17003 non-null	int64
3	ICU	17003 non-null	int64
4	MICU	17003 non-null	int64
5	NSICU	17003 non-null	int64
6	OTHER_ICU	17003 non-null	int64
7	SICU	17003 non-null	int64
8	TSICU	17003 non-null	int64
9	Bicarbonate	17003 non-null	float64
10	Calculated Bicarbonate, Whole Blood	17003 non-null	float64
11	Lactate	17003 non-null	float64
12	Oxygen	17003 non-null	float64
13	Oxygen Saturation	17003 non-null	float64
14	pCO2	17003 non-null	float64
15	pH	17003 non-null	float64
16	age_Adult	17003 non-null	int64
17	age_Senior	17003 non-null	int64
18	age_Young	17003 non-null	int64
19	admission_type_DIRECT EMER.	17003 non-null	int64
20	admission_type_DIRECT OBSERVATION	17003 non-null	int64
21	admission_type_ELECTIVE	17003 non-null	int64
22	admission_type_EU OBSERVATION	17003 non-null	int64
23	admission_type_EW EMER.	17003 non-null	int64
24	admission_type_OBSERVATION ADMIT	17003 non-null	int64
25	admission_type_SURGICAL SAME DAY ADMISSION	17003 non-null	int64
26	admission_type_URGENT	17003 non-null	int64
27	insurance_Medicaid	17003 non-null	int64
28	insurance_Medicare	17003 non-null	int64
29	insurance_No charge	17003 non-null	int64
30	insurance_Other	17003 non-null	int64
31	insurance_Private	17003 non-null	int64
32	insurance_Unknown	17003 non-null	int64
33	race_AMERICAN INDIAN/ALASKA NATIVE	17003 non-null	int64
34	race_ASIAN	17003 non-null	int64
35	race_BLACK/AFRICAN AMERICAN	17003 non-null	int64
36	race_HISPANIC/LATINO	17003 non-null	int64
37	race_NATIVE HAWAIIAN OR OTHER PACIFIC ISLANDER	17003 non-null	int64
38	race_OTHER/UNKNOWN	17003 non-null	int64
39	race_PORTUGUESE	17003 non-null	int64
40	race_SOUTH AMERICAN	17003 non-null	int64
41	race_WHITE	17003 non-null	int64
42	gender_Female	17003 non-null	int64
43	gender_Male	17003 non-null	int64
44	loc_AMBULATORY SURGERY TRANSFER	17003 non-null	int64
45	loc_CLINIC REFERRAL	17003 non-null	int64
46	loc_EMERGENCY ROOM	17003 non-null	int64
47	loc_INFORMATION NOT AVAILABLE	17003 non-null	int64
48	loc_INTERNAL TRANSFER TO OR FROM PSYCH	17003 non-null	int64
49	loc_PACU	17003 non-null	int64
50	loc_PHYSICIAN REFERRAL	17003 non-null	int64

```

51 loc_PROCEDURE SITE 17003 non-null int64
52 loc_TRANSFER FROM HOSPITAL 17003 non-null int64
53 loc_TRANSFER FROM SKILLED NURSING FACILITY 17003 non-null int64
54 loc_WALK-IN/SELF REFERRAL 17003 non-null int64
55 marital_status_DIVORCED 17003 non-null int64
56 marital_status_MARRIED 17003 non-null int64
57 marital_status_SINGLE 17003 non-null int64
58 marital_status_Unknown 17003 non-null int64
59 marital_status_WIDOWED 17003 non-null int64
dtypes: float64(7), int64(53)
memory usage: 7.8 MB

```

Splitting the Data into Training and Test Sets

```

In [ ]: # Check the distribution of hospital mortality outcomes
processed_data['hospital_expire_flag'].value_counts()

```

```

Out[ ]:

```

	count
hospital_expire_flag	
0	12424
1	4579

dtype: int64

```

In [ ]: # Create a copy of the processed data
df = processed_data.copy()
print("Original dataset size:", len(df))
print(df['hospital_expire_flag'].value_counts())

# Define features (X) and target (y)
X = df.drop(columns=['hospital_expire_flag']) # Features
y = df['hospital_expire_flag'] # Target

# Split the dataset into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, stratify=y)
print("\nTraining set size:", len(X_train))
print("Test set size:", len(X_test) ,"\n")

print('-----')
# Apply SMOTE to oversample the minority class in the training set
smote = SMOTE(sampling_strategy='auto', random_state=42)
X_train_resampled, y_train_resampled = smote.fit_resample(X_train, y_train)

print("Training set size after SMOTE:", len(X_train_resampled), "\n")

# Check class distribution after SMOTE
print(pd.Series(y_train_resampled).value_counts())

```

```
Original dataset size: 17003
hospital_expire_flag
0    12424
1     4579
Name: count, dtype: int64
```

```
Training set size: 13602
Test set size: 3401
```

```
-----
Training set size after SMOTE: 19878
```

```
hospital_expire_flag
1     9939
0     9939
Name: count, dtype: int64
```

Model Evaluation and Comparison

```
In [ ]: """
        Common function to prints model's performance metrics.
        """
def print_model_performance_metrics(name, accuracy, auc_roc, precision, recall, f1,
        """
        Prints model's performance metrics.

        Parameters:
        name (str): Name of the model.
        accuracy (float): Accuracy of the model.
        auc_roc (float): AUC-ROC of the model.
        precision (float): Precision of the model.
        recall (float): Recall of the model.
        f1 (float): F1-Score of the model.
        classification_report_output (str): Classification report of the model.
        confusion_matrix_output (ndarray): Confusion matrix of the model.
        """
    # Print performance metrics
    print(f"\n{name} Performance:")
    print(f" Accuracy: {accuracy:.4f}")
    print(f" AUC-ROC: {auc_roc:.4f}")
    print(f" Precision: {precision:.4f}")
    print(f" Recall: {recall:.4f}")
    print(f" F1-Score: {f1:.4f}")

    # Print the classification report
    print("Classification Report:")
    print(classification_report_output)

    # Print the confusion matrix
    print(f"Confusion Matrix for {name}:\n {confusion_matrix_output}")
```

Evaluate Classification Models

```

In [ ]: """
    Evaluate Classification Models
    """

    warnings.filterwarnings('ignore')

    # Standardize the features (important for neural networks)
    scaler = StandardScaler()
    X_train_resampled = scaler.fit_transform(X_train_resampled)
    X_test = scaler.transform(X_test)

    # Initialize Models
    models = {
        "Logistic Regression": LogisticRegression(random_state=0),
        "Decision Tree" : DecisionTreeClassifier(),
        "Random Forest": RandomForestClassifier(),
        "Gradient Boosting": GradientBoostingClassifier(),
        "XGBoost": XGBClassifier(learning_rate=0.1, objective='binary:logistic', random
    }

    # Prepare lists to store metrics
    metrics = []

    # Train and evaluate models on balanced data
    for name, model in models.items():
        model.fit(X_train_resampled, y_train_resampled)
        y_pred = model.predict(X_test)

        # Evaluate Model
        accuracy = accuracy_score(y_test, y_pred)
        auc_roc = roc_auc_score(y_test, model.predict_proba(X_test)[: , 1])
        precision = precision_score(y_test, y_pred)
        recall = recall_score(y_test, y_pred)
        f1 = f1_score(y_test, y_pred)

        # Append metrics for comparison
        metrics.append([accuracy, auc_roc, precision, recall, f1])

        # Print Model Performance Metrics
        cf = classification_report(y_test, y_pred)
        cm = confusion_matrix(y_test, y_pred)
        print_model_performance_metrics(name, accuracy, auc_roc, precision, recall, f1,

    # Create a DataFrame for model performance comparison
    metrics_df = pd.DataFrame(metrics, columns=['Accuracy', 'AUC-ROC', 'Precision', 'Re
    print("\nModel Performance Comparison:")
    display(metrics_df)

```

Logistic Regression Performance:

Accuracy: 0.7754

AUC-ROC: 0.7702

Precision: 0.6590

Recall: 0.3439

F1-Score: 0.4519

Classification Report:

	precision	recall	f1-score	support
0	0.79	0.93	0.86	2485
1	0.66	0.34	0.45	916
accuracy			0.78	3401
macro avg	0.73	0.64	0.66	3401
weighted avg	0.76	0.78	0.75	3401

Confusion Matrix for Logistic Regression:

[[2322 163]

[601 315]]

Decision Tree Performance:

Accuracy: 0.6998

AUC-ROC: 0.6233

Precision: 0.4443

Recall: 0.4574

F1-Score: 0.4508

Classification Report:

	precision	recall	f1-score	support
0	0.80	0.79	0.79	2485
1	0.44	0.46	0.45	916
accuracy			0.70	3401
macro avg	0.62	0.62	0.62	3401
weighted avg	0.70	0.70	0.70	3401

Confusion Matrix for Decision Tree:

[[1961 524]

[497 419]]

Random Forest Performance:

Accuracy: 0.7945

AUC-ROC: 0.7828

Precision: 0.6955

Recall: 0.4214

F1-Score: 0.5248

Classification Report:

	precision	recall	f1-score	support
0	0.81	0.93	0.87	2485
1	0.70	0.42	0.52	916
accuracy			0.79	3401
macro avg	0.75	0.68	0.70	3401
weighted avg	0.78	0.79	0.78	3401

Confusion Matrix for Random Forest:

```
[[2316 169]
 [ 530 386]]
```

Gradient Boosting Performance:

Accuracy: 0.7974
AUC-ROC: 0.7950
Precision: 0.6988
Recall: 0.4356
F1-Score: 0.5367

Classification Report:

	precision	recall	f1-score	support
0	0.82	0.93	0.87	2485
1	0.70	0.44	0.54	916
accuracy			0.80	3401
macro avg	0.76	0.68	0.70	3401
weighted avg	0.79	0.80	0.78	3401

Confusion Matrix for Gradient Boosting:

```
[[2313 172]
 [ 517 399]]
```

XGBoost Performance:

Accuracy: 0.8077
AUC-ROC: 0.8118
Precision: 0.7481
Recall: 0.4312
F1-Score: 0.5471

Classification Report:

	precision	recall	f1-score	support
0	0.82	0.95	0.88	2485
1	0.75	0.43	0.55	916
accuracy			0.81	3401
macro avg	0.78	0.69	0.71	3401
weighted avg	0.80	0.81	0.79	3401

Confusion Matrix for XGBoost:


```
[[2352 133]
 [ 521 395]]
```

Model Performance Comparison:

	Accuracy	AUC-ROC	Precision	Recall	F1-Score
Logistic Regression	0.775360	0.770217	0.658996	0.343886	0.451937
Decision Tree	0.699794	0.623279	0.444327	0.457424	0.450780
Random Forest	0.794472	0.782850	0.695495	0.421397	0.524813
Gradient Boosting	0.797413	0.794970	0.698774	0.435590	0.536651
XGBoost	0.807704	0.811770	0.748106	0.431223	0.547091

```
In [ ]: """
Common function to plot model performance metrics comparison
"""
def plot_model_metrics_comparison(metrics):
    # Plot comparison of models in a single bar plot
    metrics.plot(kind='bar', figsize=(10, 6), colormap='cividis', width=0.8)
    plt.title('Comparison of Models Performance')
    plt.ylabel('Score')
    plt.xlabel('Model')
    plt.xticks(rotation=45)
    plt.tight_layout()
    plt.show()

plot_model_metrics_comparison(metrics_df)
```



```
In [ ]: """
Plot ROC Curve comparision for all models
"""
from sklearn.metrics import roc_curve, auc
plt.figure(figsize=(8, 6))
```

```

# Plot ROC curve for each model
for name, model in models.items():
    y_proba = model.predict_proba(X_test)[: , 1]
    fpr, tpr, _ = roc_curve(y_test, y_proba)
    plt.plot(fpr, tpr, label=f"{name} (AUC = {auc(fpr, tpr):.2f})")


# Plot the diagonal line representing random classifier performance
plt.plot([0, 1], [0, 1], linestyle="--", color="gray")

# Add Labels and title
plt.xlabel("False Positive Rate", fontsize=12)
plt.ylabel("True Positive Rate", fontsize=12)
plt.title("ROC Curve Comparison", fontsize=14, fontweight='bold')

# Show the Legend
plt.legend(loc="lower right")

# Show the plot
plt.show()

```



```

In [ ]: import math

"""
Plotting Feature Importance for Model Comparison
"""

# Function to plot feature importance for each model

```

```

def plot_feature_importance(models, X_train, feature_names):
    # Dynamically calculate number of rows and columns based on the number of model
    valid_models = {name: model for name, model in models.items() if hasattr(model,
    num_models = len(valid_models)

    if num_models == 0:
        print("No models with feature importance found.")
        return

    rows = math.ceil(num_models / 3) # 3 columns per row
    cols = min(3, num_models) # Ensure we have at most 3 columns per row

    plt.figure(figsize=(16, 4 * rows)) # Adjust height based on rows

    # Iterate over models to plot feature importance
    for idx, (name, model) in enumerate(valid_models.items()):
        # For models that have feature importances
        if hasattr(model, 'feature_importances_'):
            feature_importance = model.feature_importances_
        elif hasattr(model, 'get_feature_importance'): # For models like CatBoost
            feature_importance = model.get_feature_importance()

        # Create a DataFrame for feature importances and sort it
        feature_importance_df = pd.DataFrame({
            'Feature': feature_names,
            'Importance': feature_importance
        })

        # Plot top important features
        feature_importance_df = feature_importance_df.sort_values(by='Importance',

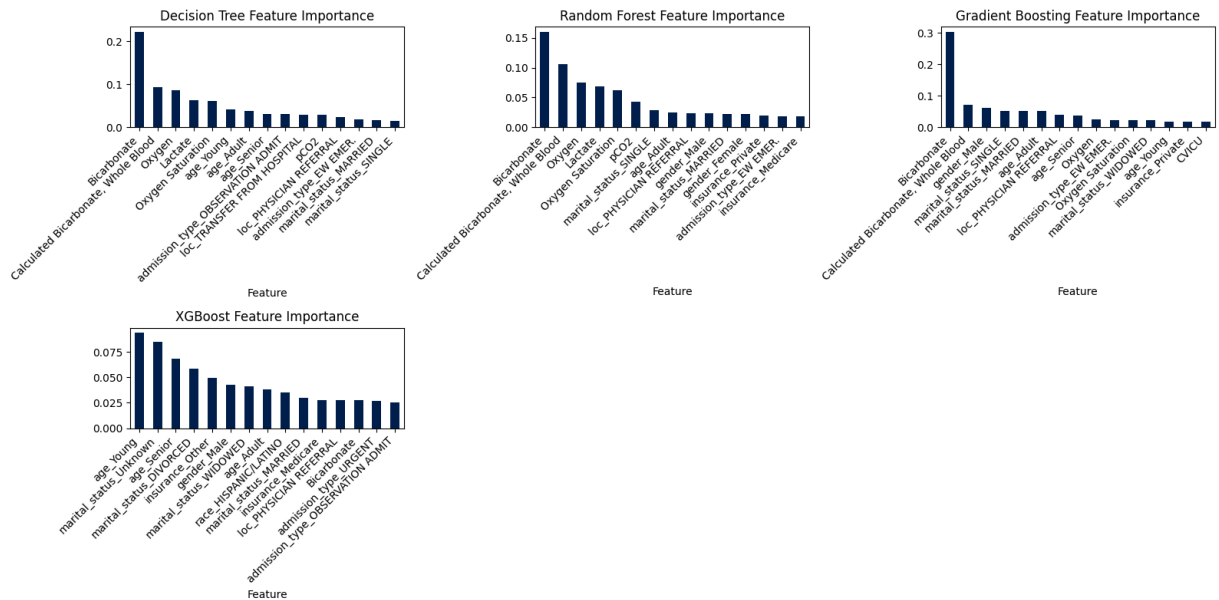
        # Define position in the grid for subplots (idx + 1 will handle 1-based ind
        ax = plt.subplot(rows, cols, idx + 1)

        # Plot feature importance for the current model
        feature_importance_df.plot.bar(x='Feature', y='Importance', legend=False, t
        plt.xticks(rotation=45, ha='right')

    plt.tight_layout()
    plt.show()

# Assuming X_train_resampled and models are defined
feature_names = X_train.columns
plot_feature_importance(models, X_train_resampled, feature_names)

```



```
In [ ]: """
Tune best model XGBoost
"""

# Define hyperparameters to tune for XGBClassifier
param_grid = {
    "n_estimators": [100, 200, 300],
    "learning_rate": [0.01, 0.1, 0.2],
    "max_depth": [3, 5, 7]
}

# Initialize model
xgb = XGBClassifier(learning_rate=0.1, objective='binary:logistic', random_state=0,

# Grid Search with 5-Fold Cross Validation
grid_search = GridSearchCV(xgb, param_grid, cv=5, scoring="roc_auc", n_jobs=-1)
grid_search.fit(X_train, y_train)

# Best parameters & best score
print(f"Best Parameters: {grid_search.best_params_}")
print(f"Best AUC-ROC Score: {grid_search.best_score_:.4f}")

# Evaluate on test data
best_xgb = grid_search.best_estimator_
y_pred_best = best_xgb.predict(X_test)

# Evaluate the best XGBoost model
accuracy = accuracy_score(y_test, y_pred_best)
auc_roc = roc_auc_score(y_test, best_xgb.predict_proba(X_test)[:, 1])
precision = precision_score(y_test, y_pred_best)
recall = recall_score(y_test, y_pred_best)
f1 = f1_score(y_test, y_pred_best)

# Append metrics for comparison
new_row = pd.Series([accuracy, auc_roc, precision, recall, f1],
                    index=metrics_df.columns, name="Tuned XGBoost")

# Use pd.concat to add the new row to the DataFrame
```

```
metrics_df = pd.concat([metrics_df, new_row.to_frame().T])

# Print Model Performance Metrics
cf = classification_report(y_test, y_pred_best)
cm = confusion_matrix(y_test, y_pred_best)
print_model_performance_metrics('XGBoost', accuracy, auc_roc, precision, recall, f1
```

Best Parameters: {'learning_rate': 0.1, 'max_depth': 3, 'n_estimators': 300}

Best AUC-ROC Score: 0.8148

XGBoost Performance:

Accuracy: 0.7236

AUC-ROC: 0.5423

Precision: 0.3696

Recall: 0.0371

F1-Score: 0.0675

Classification Report:

	precision	recall	f1-score	support
0	0.73	0.98	0.84	2485
1	0.37	0.04	0.07	916
accuracy			0.72	3401
macro avg	0.55	0.51	0.45	3401
weighted avg	0.64	0.72	0.63	3401

Confusion Matrix for XGBoost:

```
[[2427  58]
 [ 882  34]]
```

Evaluate Neural Network Model

```
In [ ]: # Standardize the features (important for neural networks)
scaler = StandardScaler()
X_train_resampled = scaler.fit_transform(X_train_resampled)
X_test = scaler.transform(X_test)

# Convert the data to PyTorch tensors
X_train_tensor = torch.tensor(X_train_resampled, dtype=torch.float32)
X_test_tensor = torch.tensor(X_test, dtype=torch.float32)
y_train_tensor = torch.tensor(y_train_resampled.values, dtype=torch.long)
y_test_tensor = torch.tensor(y_test.values, dtype=torch.long)

# Define the Deep Learning model
class ARFModel(nn.Module):
    def __init__(self, input_dim):
        super(ARFModel, self).__init__()
        self.layer11 = nn.Linear(input_dim, 128)
        self.batchnorm11 = nn.BatchNorm1d(128)
        self.layer1 = nn.Linear(128, 64)
        self.batchnorm1 = nn.BatchNorm1d(64)
        self.layer2 = nn.Linear(64, 32)
        self.batchnorm2 = nn.BatchNorm1d(32)
        self.layer3 = nn.Linear(32, 16)
        self.batchnorm3 = nn.BatchNorm1d(16) # Batch normalization
```

```

        self.output = nn.Linear(16, 2)
        self.dropout = nn.Dropout(0.3)           # Dropout layer to reduce overfitti

    def forward(self, x):
        x = F.relu(self.batchnorm11(self.layer11(x)))
        x = self.dropout(x)
        x = F.relu(self.batchnorm1(self.layer1(x)))
        x = self.dropout(x)
        x = F.relu(self.batchnorm2(self.layer2(x)))
        x = self.dropout(x)
        x = F.relu(self.batchnorm3(self.layer3(x)))
        x = self.dropout(x)
        x = self.output(x)
        return x

# Initialize model, loss function, and optimizer
input_dim = X_train_tensor.shape[1]
model = ARFModel(input_dim=input_dim)

# Compute class weights to handle imbalance in the dataset
class_weights = compute_class_weight('balanced', classes=np.array([0, 1]), y=y_train)
class_weights = torch.tensor(class_weights, dtype=torch.float32)

# Define the Loss function (CrossEntropyLoss) with class weights
criterion = nn.CrossEntropyLoss(weight=class_weights)
optimizer = optim.AdamW(model.parameters(), lr=0.001)

# Training Loop (200 epochs)
num_epochs = 200
for epoch in range(num_epochs):
    model.train()
    optimizer.zero_grad()
    outputs = model(X_train_tensor)
    loss = criterion(outputs, y_train_tensor)
    loss.backward()
    optimizer.step()

    # Print the Loss every 10 epochs
    if (epoch + 1) % 10 == 0:
        print(f"Epoch [{epoch+1}/{num_epochs}], Loss: {loss.item():.4f}")

# Evaluate the model on the test set
model.eval()
with torch.no_grad():
    outputs = model(X_test_tensor)
    _, predicted = torch.max(outputs, 1)

# Calculate various evaluation metrics
accuracy = accuracy_score(y_test_tensor, predicted)
y_prob = torch.softmax(outputs, dim=1)[ :, 1]
roc_auc = roc_auc_score(y_test_tensor, y_prob)
precision = precision_score(y_test_tensor, predicted)
recall = recall_score(y_test_tensor, predicted)
f1 = f1_score(y_test_tensor, predicted)

# Append metrics for comparison

```

```

new_row = pd.Series([accuracy, auc_roc, precision, recall, f1],
                    index=metrics_df.columns, name='Neural Network')
# Use pd.concat to add the new row to the DataFrame
metrics_df = pd.concat([metrics_df, new_row.to_frame().T])

# Print Model Performance Metrics
cf = classification_report(y_test_tensor, predicted)
cm = confusion_matrix(y_test_tensor, predicted)
print_model_performance_metrics('Neural Network', accuracy, auc_roc, precision, rec

```

```

Epoch [10/200], Loss: 0.6178
Epoch [20/200], Loss: 0.5354
Epoch [30/200], Loss: 0.4801
Epoch [40/200], Loss: 0.4466
Epoch [50/200], Loss: 0.4287
Epoch [60/200], Loss: 0.4089
Epoch [70/200], Loss: 0.3967
Epoch [80/200], Loss: 0.3861
Epoch [90/200], Loss: 0.3775
Epoch [100/200], Loss: 0.3701
Epoch [110/200], Loss: 0.3606
Epoch [120/200], Loss: 0.3559
Epoch [130/200], Loss: 0.3517
Epoch [140/200], Loss: 0.3504
Epoch [150/200], Loss: 0.3465
Epoch [160/200], Loss: 0.3384
Epoch [170/200], Loss: 0.3395
Epoch [180/200], Loss: 0.3363
Epoch [190/200], Loss: 0.3371
Epoch [200/200], Loss: 0.3305

```

Neural Network Performance:

```

Accuracy: 0.7889
AUC-ROC: 0.5423
Precision: 0.7271
Recall: 0.3461
F1-Score: 0.4689

```

Classification Report:

	precision	recall	f1-score	support
0	0.80	0.95	0.87	2485
1	0.73	0.35	0.47	916
accuracy			0.79	3401
macro avg	0.76	0.65	0.67	3401
weighted avg	0.78	0.79	0.76	3401

Confusion Matrix for Neural Network:

```

[[2366 119]
 [ 599 317]]

```

Model Comparision

```

In [ ]: #Show comparision of all the models
plot_model_metrics_comparison(metrics_df)

```

