

ML/DL Assignment

MIMIC-IV : Mortality Prediction for Acute Respiratory Failure Patients

Ankita Savaliya

MIMIC-IV : Mortality Prediction for Acute Respiratory Failure Patients

Goal:

- Load and preprocess MIMIC-IV EHR data related to Acute Respiratory Failure (ARF).
- Analyze demographic and lab test indicators for ARF.
- Construct a structured dataset for modeling.
- Apply classification and deep learning models to EHR data.
- Compare model performances.
- Predict mortality in ARF patients.

Dataset: [MIMIC-IV](#) dataset.

GitHub and Google Colab Links:

https://colab.research.google.com/github/AnkitaSavaliya/AIH/blob/main/MIMIC-IV_MORTALITY_PREDICTION_ARF.ipynb

https://github.com/AnkitaSavaliya/AIH/blob/main/MIMIC-IV_MORTALITY_PREDICTION_ARF.ipynb

https://github.com/AnkitaSavaliya/AIH/blob/main/ML-DL_MIMIC-IV_ARF_MORTALITY_PREDICTION.pptx

Step 1: Import required libraries

```
import pandas as pd
import numpy as np
import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, roc_auc_score, precision_score, recall_score, f1_score, classification_report, confusion_matrix
from imblearn.over_sampling import SMOTE
from sklearn.utils.class_weight import compute_class_weight
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from xgboost import XGBClassifier
import warnings

# Mount Google Drive
from google.colab import drive
drive.mount('/content/drive')
```

Import libraries in Google Colab like pandas, sklearn, matplotlib, torch etc.
Mount Google Drive where MIMIC-IV CSV files are stored.

Step 2: Loading and Filter Data for ARF

```
def read_mimic_csv_file(mimic_csv_file_name: str, low_memory: bool = False, chunksize: int = None) -> pd.DataFrame:
    """
    Read a CSV file from the MIMIC-IV dataset into a pandas DataFrame.

    Parameters:
    - mimic_csv_file_name (str): Name of the CSV file.
    - low_memory (bool): Whether to use low memory mode when reading.
    - chunksize (int, optional): Number of rows per chunk if reading in chunks.

    Returns:
    - pd.DataFrame
    """
    # Define the root directory of MIMIC-IV data in Google Drive
    mimic_root_dir_path = "/content/drive/MyDrive/Colab Notebooks/AIH/MIMIC-IV/"
    file_path = mimic_root_dir_path + mimic_csv_file_name

    return pd.read_csv(file_path, low_memory=low_memory, chunksize=chunksize)
```

This function reads a given CSV file from the MIMIC-IV dataset (stored in Google Drive) and returns a pandas DataFrame. Some files, like *labevents*, are very large, so this method supports reading in chunks to optimize memory usage.

Step 2: Loading and Filter Data for ARF(Continued)...

This part filters diagnoses using ICD-9 and ICD-10 codes related to Acute Respiratory Failure (ARF) and merges the filtered data with the admissions, patients, and ICU stays tables. Additionally, unnecessary columns are dropped, duplicates are removed, and the dataset is reset for a clean structure.

```
arf_diagnoses_df = read_mimic_csv_file("diagnoses_icd.csv.gz")

# Define relevant ICD-9 and ICD-10 codes for acute respiratory failure(MIMIC-IV contains both ICD-9 and ICD-10 codes)
arf_icd_codes = {'51851', '51881', 'J960', 'J9600', 'J9601', 'J9602'}

# Filter diagnoses dataset
arf_diagnoses_df = arf_diagnoses_df[arf_diagnoses_df['icd_code'].isin(arf_icd_codes)].copy()

# Drop unnecessary columns
arf_diagnoses_df.drop(columns=['seq_num', 'icd_code', 'icd_version'], inplace=True, errors='ignore')

# Remove duplicates
arf_diagnoses_df.drop_duplicates(inplace=True)

# Merge with admissions data
arf_admissions_df = read_mimic_csv_file('admissions.csv.gz')

arf_merged_df = arf_diagnoses_df.merge(
    arf_admissions_df, on=['subject_id', 'hadm_id'], how='inner'
)

arf_merged_df.drop(columns=['dischtime', 'deathtime', 'admit_provider_id', 'discharge_location',
                             'language', 'edregtime', 'edouttime'], inplace=True, errors='ignore')

arf_merged_df.drop_duplicates(inplace=True)
arf_merged_df.reset_index(drop=True, inplace=True)

# Merge with patient demographics
arf_patients_df = read_mimic_csv_file('patients.csv.gz')

arf_merged_df = arf_merged_df.merge(
    arf_patients_df, on=['subject_id'], how='inner'
)

arf_merged_df.drop(columns=['dod', 'anchor_year_group'], inplace=True, errors='ignore')
arf_merged_df.drop_duplicates(inplace=True)
arf_merged_df.reset_index(drop=True, inplace=True)

# Merge with ICU stays
arf_icustays_df = read_mimic_csv_file('icustays.csv.gz')

arf_merged_df = arf_merged_df.merge(
    arf_icustays_df, on=['subject_id', 'hadm_id'], how='inner'
)
```

Step 2: Loading and Filter Data for ARF(Continued)...

```
arf_merged_df.drop(columns=['last_careunit', 'intime', 'outtime', 'los', 'stay_id'], inplace=True, errors='ignore')
arf_merged_df.drop_duplicates(inplace=True)
arf_merged_df.reset_index(drop=True, inplace=True)

# Define lab test keywords related to respiratory function
resp_lab_tests = {
    'oxygen saturation', 'oxygen', 'ph', 'pco2',
    'bicarbonate', 'lactate', 'calculated bicarbonate, whole blood'
}

# Load lab item details
lab_items_df = read_mimic_csv_file('d_labitems.csv.gz')

# Filter respiratory-related blood lab items
lab_items_df = lab_items_df[
    (lab_items_df['fluid'] == 'Blood') &
    (lab_items_df['label'].str.lower().str.strip().isin(resp_lab_tests))
].copy()

# Drop unnecessary columns
lab_items_df.drop(columns=['fluid', 'category'], inplace=True, errors='ignore')
lab_items_df.drop_duplicates(inplace=True)
lab_items_df.reset_index(drop=True, inplace=True)

# Extract unique subject_id and hadm_id pairs
subject_hadm_set = arf_merged_df[['subject_id', 'hadm_id']].drop_duplicates().reset_index(drop=True)
```

- Lab tests, such as oxygen saturation, pH, pCO₂, and bicarbonate levels, are critical indicators of respiratory function and can serve as important features for mortality prediction.
- This code filters ARF-specific lab events and merges the datasets.
- To manage the large volume of lab event data efficiently, we processed it in chunks, ensuring memory optimization throughout the process.

```
# Process lab events data in chunks to manage memory efficiently
lab_chunks = []
for lab_chunk in read_mimic_csv_file('labevents.csv.gz', low_memory=False, chunksize=10**7):
    # Drop irrelevant columns
    lab_chunk.drop(columns=['labevent_id', 'value', 'valueuom', 'flag', 'ref_range_lower', 'ref_range_upper',
                           'priority', 'specimen_id', 'order_provider_id', 'storetime', 'comments'],
                  inplace=True, errors='ignore')

    # Merge with filtered lab items
    lab_chunk = lab_chunk.merge(lab_items_df, on='itemid', how='inner')
    lab_chunk.drop(columns=['itemid'], inplace=True, errors='ignore')

    # Keep only data for acute respiratory failure patients
    lab_chunk = lab_chunk.merge(subject_hadm_set, on=['subject_id', 'hadm_id'], how='inner')

    # Sort for time-based aggregation
    lab_chunk.sort_values(by=['subject_id', 'hadm_id', 'charttime'], inplace=True)

    # Aggregate lab test values by median per subject_id, hadm_id, and label
    lab_chunk = lab_chunk.groupby(['subject_id', 'hadm_id', 'label'], as_index=False)['valuenum'].median()

    lab_chunks.append(lab_chunk)

# Merge processed lab event data with the main dataset
if lab_chunks:
    arf_merged_df = arf_merged_df.merge(pd.concat(lab_chunks, ignore_index=True),
                                       on=['subject_id', 'hadm_id'], how='inner')

# Remove duplicate rows
arf_merged_df.drop_duplicates(subset=['subject_id', 'hadm_id', 'label'], inplace=True)
arf_merged_df.reset_index(drop=True, inplace=True)
```

Step 2: Loading and Filter Data for ARF(Continued)...

- After loading and filtering the data, the columns include patient identifiers, admission details, lab test results, and demographic information related to ARF.
- The next step involves understanding, processing, and cleaning the data to ensure it's ready for modeling. This may include handling missing values, transforming variables, and ensuring consistency across datasets.

```
# Display dataset info
```

```
arf_merged_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 83237 entries, 0 to 83236
```

```
Data columns (total 15 columns):
```

#	Column	Non-Null Count	Dtype
0	subject_id	83237 non-null	int64
1	hadm_id	83237 non-null	int64
2	admittime	83237 non-null	object
3	admission_type	83237 non-null	object
4	admission_location	83237 non-null	object
5	insurance	81768 non-null	object
6	marital_status	72965 non-null	object
7	race	83237 non-null	object
8	hospital_expire_flag	83237 non-null	int64
9	gender	83237 non-null	object
10	anchor_age	83237 non-null	int64
11	anchor_year	83237 non-null	int64
12	first_careunit	83237 non-null	object
13	label	83237 non-null	object
14	valuenum	83213 non-null	float64

```
dtypes: float64(1), int64(5), object(9)
```

```
memory usage: 9.5+ MB
```

Step 3: Preprocessing and Feature Engineering for ARF

This step performs preprocessing (cleaning, mapping) and feature engineering on the dataset created earlier. The following processing is done:

- Mapping gender, handling missing values in marital status and insurance.
- Calculating age and categorizing patients into age groups.
- Standardizing race and ICU categories and creating separate columns for each ICU unit.
- Removing unnecessary columns to prepare the dataset for analysis.

```
# Create a copy of the merged Acute Respiratory Failure dataset for processing
arf_processed_df = arf_merged_df.copy()

# Map Gender Column
arf_processed_df['gender'] = arf_processed_df['gender'].map({'F': 'Female', 'M': 'Male'})

# Handle missing values in marital status by replacing NaNs with 'Unknown'
arf_processed_df['marital_status'] = arf_processed_df['marital_status'].fillna('Unknown')

# Handle missing values in insurance by replacing NaNs with 'Unknown'
arf_processed_df['insurance'] = arf_processed_df['insurance'].fillna('Unknown')

# Convert admission time to datetime format
arf_processed_df['admittime'] = pd.to_datetime(arf_processed_df['admittime'])

# Compute patient age at admission using MIMIC-IV anchor values
arf_processed_df['admission_age'] = (
    arf_processed_df['anchor_age'] +
    (arf_processed_df['admittime'].dt.year - arf_processed_df['anchor_year'])
)

# Categorize patients into age groups: Young (<30), Adult (30-60), Senior (60+)
arf_processed_df['age_group'] = pd.cut(
    arf_processed_df['admission_age'],
    bins=[0, 30, 60, float('inf')],
    labels=['Young', 'Adult', 'Senior'],
    right=False
)

# Remove unnecessary columns after computing age group
arf_processed_df.drop(columns=['admittime', 'anchor_year', 'anchor_age', 'admission_age'], inplace=True)

# Convert age group to string type
arf_processed_df['age_group'] = arf_processed_df['age_group'].astype(str)
```

```
# Standardize race categories by grouping similar values
arf_processed_df['race'] = arf_processed_df['race'].replace(
    {r"ASIAN\D*": "ASIAN",
    r"WHITE\D*": "WHITE",
    r"HISPANIC\D*": "HISPANIC/LATINO",
    r"BLACK\D*": "BLACK/AFRICAN AMERICAN"},
    regex=True
)

# Replace ambiguous race values with 'OTHER/UNKNOWN'
arf_processed_df['race'] = arf_processed_df['race'].replace(
    ['UNABLE TO OBTAIN', 'OTHER', 'PATIENT DECLINED TO ANSWER', 'UNKNOWN', 'MULTIPLE RACE/ETHNICITY'],
    'OTHER/UNKNOWN'
)

# Standardize ICU (first care unit) categories by grouping related units
arf_processed_df['first_careunit'] = arf_processed_df['first_careunit'].replace(
    {r"Medical/Surgical\D*": "MICU, SICU",
    r"Medical\D*": "MICU",
    r"Neuro\D*": "NSICU",
    r"Cardiac\D*": "CVICU",
    r"Coronary\D*": "CCU",
    r"Trauma SICU\D*": "TSICU",
    r"Surgical\D*": "SICU",
    r"Intensive Care Unit\D*": "ICU"},
    regex=True
)

# Convert uncommon ICU categories into 'OTHERICU'
arf_processed_df['first_careunit'] = arf_processed_df['first_careunit'].replace(
    ['Surgery/Vascular/Intermediate', 'PACU', 'Medicine', 'Surgery/Trauma', 'Med/Surg', 'Neuro Stepdown'],
    'OTHER_ICU'
)

# Convert ICU categories into separate binary columns (one-hot encoding)
arf_processed_df['first_careunit'] = arf_processed_df['first_careunit'].str.split(' ', expand=False).reset_index(drop=True)
arf_processed_df = arf_processed_df.join(
    pd.get_dummies(arf_processed_df['first_careunit'].apply(pd.Series).stack(), dtype=int)
    .groupby(level=0)
    .sum()
    .how='outer'
)

# Remove the original ICU category column after encoding
arf_processed_df.drop(columns=['first_careunit'], inplace=True)
```


Step 3: Preprocessing and Feature Engineering for ARF(Continued)...

Continuing data processing:

- Aggregate lab test results for each patient encounter.
- Transform lab test names into separate columns.
- Replace missing values (NaNs) with 0.
- Convert categorical features (admission type, insurance, race, gender, admission location, and marital status) into a binary format.
- Remove duplicates and reset indices to finalize the dataset for model training.

```
# Aggregate lab test results by subject_id and hadm_id
tmp = arf_processed_df.groupby(['subject_id', 'hadm_id'], as_index=False)[['label', 'valuenum']].agg(list).reset_index(drop=True)

# Drop old lab event columns since they have been aggregated
arf_processed_df.drop(columns=['label', 'valuenum'], inplace=True)

# Merge aggregated lab results back into the main dataframe
arf_processed_df = arf_processed_df.merge(tmp, on=['subject_id', 'hadm_id'], how='inner')

# Clean up temporary variable
del tmp

# Extract unique lab test names from the 'label' column
all_labels = sorted(set(itertools.chain.from_iterable(arf_processed_df['label'])))

# Expand 'valuenum' into separate columns with lab test names as headers
arf_processed_df = arf_processed_df.join(
    pd.DataFrame(arf_processed_df['valuenum'].to_list(), columns=all_labels,
        how="outer"
    )
)

# Drop unnecessary columns after transformation
arf_processed_df.drop(columns=['subject_id', 'hadm_id', 'label', 'valuenum'], inplace=True, errors='ignore')

# Handle missing values by replacing NaNs with 0
arf_processed_df.fillna(0, inplace=True)

# One-hot encode category columns: admission type, insurance, race, gender, admission location, and marital status
prefix_cols = ['age', "admission_type", "insurance", 'race', 'gender', 'loc', 'marital_status']
dummy_cols = ['age_group', 'admission_type', 'insurance', 'race', 'gender', 'admission_location', 'marital_status']
arf_processed_df = pd.get_dummies(arf_processed_df, prefix=prefix_cols, columns=dummy_cols, dtype=int)

# Drop duplicates, drop rows with NaN, and reset indices
arf_processed_df.drop_duplicates(inplace=True)
arf_processed_df.dropna(inplace=True)
arf_processed_df.reset_index(drop=True, inplace=True)
```

Step 3: Preprocessing and Feature Engineering for ARF(Continued)...

After processing and feature engineering, the dataset now contains the following columns. It is now ready for model training.

```
# Exploring the columns of preprocessed data
processed_data = arf_processed_df.copy()
processed_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 17003 entries, 0 to 17002
Data columns (total 60 columns):
```

#	Column	Non-Null Count	Dtype
0	hospital_expire_flag	17003 non-null	int64
1	CCU	17003 non-null	int64
2	CVICU	17003 non-null	int64
3	ICU	17003 non-null	int64
4	MICU	17003 non-null	int64
5	NSICU	17003 non-null	int64
6	OTHER_ICU	17003 non-null	int64
7	SICU	17003 non-null	int64
8	TSICU	17003 non-null	int64
9	Bicarbonate	17003 non-null	float64
10	Calculated Bicarbonate, Whole Blood	17003 non-null	float64
11	Lactate	17003 non-null	float64
12	Oxygen	17003 non-null	float64
13	Oxygen Saturation	17003 non-null	float64
14	pCO2	17003 non-null	float64
15	pH	17003 non-null	float64
16	age_Adult	17003 non-null	int64
17	age_Senior	17003 non-null	int64
18	age_Young	17003 non-null	int64
19	admission_type_DIRECT EMER.	17003 non-null	int64
20	admission_type_DIRECT OBSERVATION	17003 non-null	int64
21	admission_type_ELECTIVE	17003 non-null	int64
22	admission_type_EU OBSERVATION	17003 non-null	int64
23	admission_type_EW EMER.	17003 non-null	int64
24	admission_type_OBSERVATION ADMIT	17003 non-null	int64
25	admission_type_SURGICAL SAME DAY ADMISSION	17003 non-null	int64
26	admission_type_URGENT	17003 non-null	int64
27	insurance_Medicaid	17003 non-null	int64
28	insurance_Medicare	17003 non-null	int64

29	insurance_No charge	17003 non-null	int64
30	insurance_Other	17003 non-null	int64
31	insurance_Private	17003 non-null	int64
32	insurance_Unknown	17003 non-null	int64
33	race_AMERICAN INDIAN/ALASKA NATIVE	17003 non-null	int64
34	race_ASIAN	17003 non-null	int64
35	race_BLACK/AFRICAN AMERICAN	17003 non-null	int64
36	race_HISPANIC/LATINO	17003 non-null	int64
37	race_NATIVE HAWAIIAN OR OTHER PACIFIC ISLANDER	17003 non-null	int64
38	race_OTHER/UNKNOWN	17003 non-null	int64
39	race_PORTUGUESE	17003 non-null	int64
40	race_SOUTH AMERICAN	17003 non-null	int64
41	race_WHITE	17003 non-null	int64
42	gender_Female	17003 non-null	int64
43	gender_Male	17003 non-null	int64
44	loc_AMBULATORY SURGERY TRANSFER	17003 non-null	int64
45	loc_CLINIC REFERRAL	17003 non-null	int64
46	loc_EMERGENCY ROOM	17003 non-null	int64
47	loc_INFORMATION NOT AVAILABLE	17003 non-null	int64
48	loc_INTERNAL TRANSFER TO OR FROM PSYCH	17003 non-null	int64
49	loc_PACU	17003 non-null	int64
50	loc_PHYSICIAN REFERRAL	17003 non-null	int64
51	loc_PROCEDURE SITE	17003 non-null	int64
52	loc_TRANSFER FROM HOSPITAL	17003 non-null	int64
53	loc_TRANSFER FROM SKILLED NURSING FACILITY	17003 non-null	int64
54	loc_WALK-IN/SELF REFERRAL	17003 non-null	int64
55	marital_status_DIVORCED	17003 non-null	int64
56	marital_status_MARRIED	17003 non-null	int64
57	marital_status_SINGLE	17003 non-null	int64
58	marital_status_Unknown	17003 non-null	int64
59	marital_status_WIDOWED	17003 non-null	int64

Step 4: Splitting the Data into Training and Test Sets

This step prepares the dataset for modeling:

- Defines features (X) and target (y).
- Splits the dataset into training (80%) and test (20%) sets while preserving class distribution.
- As the data is imbalanced(12424 negative case and 4579 positive case), applies SMOTE to oversample the minority class in the training set, ensuring a balanced dataset for better model performance.

```
# Check the distribution of hospital mortality outcomes
processed_data['hospital_expire_flag'].value_counts()
```

count	
hospital_expire_flag	
0	12424
1	4579

```
# Create a copy of the processed data
df = processed_data.copy()
print("Original dataset size:", len(df))
print(df['hospital_expire_flag'].value_counts())

# Define features (X) and target (y)
X = df.drop(columns=['hospital_expire_flag']) # Features
y = df['hospital_expire_flag'] # Target

# Split the dataset into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, stratify=y, random_state=42)
print("\nTraining set size:", len(X_train))
print("Test set size:", len(X_test), "\n")

print('-----')
# Apply SMOTE to oversample the minority class in the training set
smote = SMOTE(sampling_strategy='auto', random_state=42)
X_train_resampled, y_train_resampled = smote.fit_resample(X_train, y_train)

print("Training set size after SMOTE:", len(X_train_resampled), "\n")

# Check class distribution after SMOTE
print(pd.Series(y_train_resampled).value_counts())
```

```
Original dataset size: 17003
hospital_expire_flag
0    12424
1     4579
Name: count, dtype: int64
```

```
Training set size: 13602
Test set size: 3401
```

```
-----
Training set size after SMOTE: 19878
```

```
hospital_expire_flag
1     9939
0     9939
Name: count, dtype: int64
```

Step 5: Model Evaluation and Comparison (sklearn models and XGBoost)

```
"""
    Evaluate Classification Models
    """
warnings.filterwarnings('ignore')

# Standardize the features (important for neural networks)
scaler = StandardScaler()
X_train_resampled = scaler.fit_transform(X_train_resampled)
X_test = scaler.transform(X_test)

# Initialize Models
models = {
    "Logistic Regression": LogisticRegression(random_state=0),
    "Decision Tree": DecisionTreeClassifier(),
    "Random Forest": RandomForestClassifier(),
    "Gradient Boosting": GradientBoostingClassifier(),
    "XGBoost": XGBClassifier(learning_rate=0.1, objective='binary:logistic', random_state=0, eval_metric='mlogloss')
}

# Prepare lists to store metrics
metrics = []

# Train and evaluate models on balanced data
for name, model in models.items():
    model.fit(X_train_resampled, y_train_resampled)
    y_pred = model.predict(X_test)

    # Evaluate Model
    accuracy = accuracy_score(y_test, y_pred)
    auc_roc = roc_auc_score(y_test, model.predict_proba(X_test)[:, 1])
    precision = precision_score(y_test, y_pred)
    recall = recall_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)

    # Append metrics for comparison
    metrics.append([accuracy, auc_roc, precision, recall, f1])

# Print Model Performance Metrics
cf = classification_report(y_test, y_pred)
cm = confusion_matrix(y_test, y_pred)
print_model_performance_metrics(name, accuracy, auc_roc, precision, recall, f1, cf, cm)

# Create a DataFrame for model performance comparison
metrics_df = pd.DataFrame(metrics, columns=['Accuracy', 'AUC-ROC', 'Precision', 'Recall', 'F1-Score'], index=models.keys())
print("\nModel Performance Comparison:")
display(metrics_df)
```

This step evaluates various classification models to predict mortality in ARF patients. The features are standardized to improve model performance. Multiple models—Logistic Regression, Decision Tree, Random Forest, Gradient Boosting, and XGBoost—are trained and tested. The models are evaluated and compared using the following performance metrics:

Accuracy: The proportion of correct predictions out of all predictions.

AUC-ROC: A measure of the model's ability to distinguish between classes.

Precision: The proportion of true positive predictions among all positive predictions.

Recall: The proportion of actual positive cases correctly identified by the model.

F1-Score: The harmonic mean of precision and recall, balancing both metrics.

This comparison helps identify the best-performing model for predicting mortality in ARF patients.

```
def print_model_performance_metrics(name, accuracy, auc_roc, precision, recall, f1, classification_report_output, confusion_matrix_output):
    """
    Prints model's performance metrics.

    Parameters:
    name (str): Name of the model.
    accuracy (float): Accuracy of the model.
    auc_roc (float): AUC-ROC of the model.
    precision (float): Precision of the model.
    recall (float): Recall of the model.
    f1 (float): F1-Score of the model.
    classification_report_output (str): Classification report of the model.
    confusion_matrix_output (ndarray): Confusion matrix of the model.
    """

    # Print performance metrics
    print(f"\n{name} Performance:")
    print(f" Accuracy: {accuracy:.4f}")
    print(f" AUC-ROC: {auc_roc:.4f}")
    print(f" Precision: {precision:.4f}")
    print(f" Recall: {recall:.4f}")
    print(f" F1-Score: {f1:.4f}")

    # Print the classification report
    print("Classification Report:")
    print(classification_report_output)

    # Print the confusion matrix
    print(f"Confusion Matrix for {name}:\n {confusion_matrix_output}")
```

Step 5: Model Evaluation and Comparison (sklearn models and XGBoost) (Continued)...

Following is outcome of how each model performed.

Logistic Regression Performance:

Accuracy: 0.7754
AUC-ROC: 0.7702
Precision: 0.6590
Recall: 0.3439
F1-Score: 0.4519

Classification Report:

	precision	recall	f1-score	support
0	0.79	0.93	0.86	2485
1	0.66	0.34	0.45	916
accuracy			0.78	3401
macro avg	0.73	0.64	0.66	3401
weighted avg	0.76	0.78	0.75	3401

Confusion Matrix for Logistic Regression:

```
[[2322 163]
 [ 601 315]]
```

Decision Tree Performance:

Accuracy: 0.6998
AUC-ROC: 0.6233
Precision: 0.4443
Recall: 0.4574
F1-Score: 0.4508

Classification Report:

	precision	recall	f1-score	support
0	0.80	0.79	0.79	2485
1	0.44	0.46	0.45	916
accuracy			0.70	3401
macro avg	0.62	0.62	0.62	3401
weighted avg	0.70	0.70	0.70	3401

Confusion Matrix for Decision Tree:

```
[[1961 524]
 [ 497 419]]
```

Random Forest Performance:

Accuracy: 0.7945
AUC-ROC: 0.7828
Precision: 0.6955
Recall: 0.4214
F1-Score: 0.5248

Classification Report:

	precision	recall	f1-score	support
0	0.81	0.93	0.87	2485
1	0.70	0.42	0.52	916
accuracy			0.79	3401
macro avg	0.75	0.68	0.70	3401
weighted avg	0.78	0.79	0.78	3401

Confusion Matrix for Random Forest:

```
[[2316 169]
 [ 530 386]]
```

Gradient Boosting Performance:

Accuracy: 0.7974
AUC-ROC: 0.7950
Precision: 0.6988
Recall: 0.4356
F1-Score: 0.5367

Classification Report:

	precision	recall	f1-score	support
0	0.82	0.93	0.87	2485
1	0.70	0.44	0.54	916
accuracy			0.80	3401
macro avg	0.76	0.68	0.70	3401
weighted avg	0.79	0.80	0.78	3401

Confusion Matrix for Gradient Boosting:

```
[[2313 172]
 [ 517 399]]
```

XGBoost Performance:

Accuracy: 0.8077
AUC-ROC: 0.8118
Precision: 0.7481
Recall: 0.4312
F1-Score: 0.5471

Classification Report:

	precision	recall	f1-score	support
0	0.82	0.95	0.88	2485
1	0.75	0.43	0.55	916
accuracy			0.81	3401
macro avg	0.78	0.69	0.71	3401
weighted avg	0.80	0.81	0.79	3401

Confusion Matrix for XGBoost:

```
[[2352 133]
 [ 521 395]]
```


	Accuracy	AUC-ROC	Precision	Recall	F1-Score
Logistic Regression	0.775360	0.770217	0.658996	0.343886	0.451937
Decision Tree	0.699794	0.623279	0.444327	0.457424	0.450780
Random Forest	0.794472	0.782850	0.695495	0.421397	0.524813
Gradient Boosting	0.797413	0.794970	0.698774	0.435590	0.536651
XGBoost	0.807704	0.811770	0.748106	0.431223	0.547091

Step 5: Model Evaluation and Comparison (sklearn models and XGBoost) (Continued)...

The right-side plot visualizes the model comparison, highlighting the performance differences across various metrics. As observed from the metrics and visuals:

- Logistic Regression struggles with recall (34.4%), missing many mortality cases, but maintains decent precision (65.9%).
- Decision Tree performs the worst overall, with the lowest AUC-ROC (62.3%) and precision (44.3%).
- Random Forest and Gradient Boosting show improvements in recall (42.1% and 43.5%, respectively), meaning they identify more mortality cases.
- XGBoost performs the best, achieving the highest accuracy (80.8%) and AUC-ROC (81.1%), demonstrating a better balance of precision (74.8%) and recall (43.1%) compared to the other models.
- So overall XGBoost is best performing model.

```
def plot_model_metrics_comparison(metrics):  
    # Plot comparison of models in a single bar plot  
    metrics.plot(kind='bar', figsize=(10, 6), colormap='cividis', width=0.8)  
    plt.title('Comparison of Models Performance')  
    plt.ylabel('Score')  
    plt.xlabel('Model')  
    plt.xticks(rotation=45)  
    plt.tight_layout()  
    plt.show()  
  
plot_model_metrics_comparison(metrics_df)
```



Step 5: Model Evaluation and Comparison (sklearn models and XGBoost) (Continued)...

This shows the ROC curves, which visualize the True Positive Rate (TPR) and False Positive Rate (FPR) for each classification model. By comparing the curves, we can identify which model best balances the trade-off between false positives and true positives. A higher AUC indicates better performance in distinguishing between the mortality and survival classes.

```
from sklearn.metrics import roc_curve, auc
plt.figure(figsize=(8, 6))


# Plot ROC curve for each model
for name, model in models.items():
    y_proba = model.predict_proba(X_test)[: , 1]
    fpr, tpr, _ = roc_curve(y_test, y_proba)
    plt.plot(fpr, tpr, label=f"{name} (AUC = {auc(fpr, tpr):.2f})")

# Plot the diagonal line representing random classifier performance
plt.plot([0, 1], [0, 1], linestyle="--", color="gray")

# Add labels and title
plt.xlabel("False Positive Rate", fontsize=12)
plt.ylabel("True Positive Rate", fontsize=12)
plt.title("ROC Curve Comparison", fontsize=14, fontweight='bold')

# Show the legend
plt.legend(loc="lower right")

# Show the plot
plt.show()
```



Step 5: Model Evaluation and Comparison (sklearn models and XGBoost) (Continued)...

- This code visualizes the feature importance for each model, helping us understand which features have the most influence on the model's predictions.
- Code plots the top 15 most important features for each model that supports `feature_importances_`.

```
def plot_feature_importance(models, X_train, feature_names):
    # Dynamically calculate number of rows and columns based on the number of models with feature importances
    valid_models = {name: model for name, model in models.items() if hasattr(model, 'feature_importances_') or hasattr(model, 'get_feature_importance')}
    num_models = len(valid_models)

    if num_models == 0:
        print("No models with feature importance found.")
        return

    rows = math.ceil(num_models / 3) # 3 columns per row
    cols = min(3, num_models) # Ensure we have at most 3 columns per row

    plt.figure(figsize=(16, 4 * rows)) # Adjust height based on rows

    # Iterate over models to plot feature importance
    for idx, (name, model) in enumerate(valid_models.items()):
        # For models that have feature importances
        if hasattr(model, 'feature_importances_'):
            feature_importance = model.feature_importances_
        elif hasattr(model, 'get_feature_importance'): # For models like CatBoost
            feature_importance = model.get_feature_importance()

        # Create a DataFrame for feature importances and sort it
        feature_importance_df = pd.DataFrame({
            'Feature': feature_names,
            'Importance': feature_importance
        })

        # Plot top important features
        feature_importance_df = feature_importance_df.sort_values(by='Importance', ascending=False).head(15)





        # Define position in the grid for subplots (idx + 1 will handle 1-based indexing in subplot)
        ax = plt.subplot(rows, cols, idx + 1)

        # Plot feature importance for the current model
        feature_importance_df.plot.bar(x='Feature', y='Importance', legend=False, title=f"{name} Feature Importance", ax=ax, colormap='cividis')
        plt.xticks(rotation=45, ha='right')

    plt.tight_layout()
    plt.show()

# Assuming X_train_resampled and models are defined
feature_names = X_train.columns
plot_feature_importance(models, X_train_resampled, feature_names)
```


Step 5: Model Evaluation and Comparison (sklearn models and XGBoost) (Continued)...



The feature importance plots strongly suggest that lab tests, particularly oxygen-related measures and bicarbonate levels, are crucial predictors across all the models.

Step 5: Model Evaluation and Comparison (sklearn models and XGBoost) (Continued)...

```
# Define hyperparameters to tune for XGBClassifier
param_grid = {
    "n_estimators": [100, 200, 300],
    "learning_rate": [0.01, 0.1, 0.2],
    "max_depth": [3, 5, 7]
}

# Initialize model
xgb = XGBClassifier(learning_rate=0.1, objective='binary:logistic', random_state=0, eval_metric='mlogloss')

# Grid Search with 5-Fold Cross Validation
grid_search = GridSearchCV(xgb, param_grid, cv=5, scoring="roc_auc", n_jobs=-1)
grid_search.fit(X_train, y_train)

# Best parameters & best score
print(f"Best Parameters: {grid_search.best_params_}")
print(f"Best AUC-ROC Score: {grid_search.best_score_:.4f}")

# Evaluate on test data
best_xgb = grid_search.best_estimator_
y_pred_best = best_xgb.predict(X_test)

# Evaluate the best XGBoost model
accuracy = accuracy_score(y_test, y_pred_best)
auc_roc = roc_auc_score(y_test, best_xgb.predict_proba(X_test)[:, 1])
precision = precision_score(y_test, y_pred_best)
recall = recall_score(y_test, y_pred_best)
f1 = f1_score(y_test, y_pred_best)

# Append metrics for comparison
new_row = pd.Series([accuracy, auc_roc, precision, recall, f1],
                    index=metrics_df.columns, name="Tuned XGBoost")

# Use pd.concat to add the new row to the DataFrame
metrics_df = pd.concat([metrics_df, new_row.to_frame().T])

# Print Model Performance Metrics
cf = classification_report(y_test, y_pred_best)
cm = confusion_matrix(y_test, y_pred_best)
print_model_performance_metrics('XGBoost', accuracy, auc_roc, precision, recall, f1, cf, cm)
```

```
Best Parameters: {'learning_rate': 0.1, 'max_depth': 3, 'n_estimators': 300}
Best AUC-ROC Score: 0.8148
```

```
XGBoost Performance:
Accuracy: 0.7236
AUC-ROC: 0.5423
Precision: 0.3696
Recall: 0.0371
F1-Score: 0.0675
```

```
Classification Report:
              precision    recall  f1-score   support

     0       0.73         0.98         0.84         2485
     1       0.37         0.04         0.07          916

 accuracy          0.72         0.72         0.72         3401
 macro avg         0.55         0.51         0.45         3401
 weighted avg         0.64         0.72         0.63         3401
```

```
Confusion Matrix for XGBoost:
[[2427  58]
 [ 882  34]]
```

This shows hyperparameter tuning of best performing model(XGBoost) we found in earlier steps.

- After tuning, unfortunately model's performance worsened. The model's accuracy dropped from 80.7% to 72.3%, and more importantly, its recall for mortality cases significantly declined to 3.7%.
- This suggests that the tuned hyperparameters may have overfitted to training data or altered the balance between precision and recall, making the model less effective at identifying critical cases. So our earlier model was better.

Step 6: Evaluate Neural Network Model

Here, the Neural Network model is evaluated for predicting mortality in ARF patients.

- Features are standardized to ensure that all inputs are on a similar scale, improving model stability.
- Training and test datasets are converted into PyTorch tensors for compatibility with the deep learning framework.
- Multiple fully connected layers are used to capture complex patterns.
- Batch Normalization is applied to stabilize training and improve generalization.
- Dropout layers are added to reduce overfitting.
- CrossEntropyLoss is used since the task involves classification (mortality prediction).
- AdamW optimizer is chosen for efficient weight updates.

Next, we will evaluate this model and compare its performance with earlier models.

```
# Standardize the features (important for neural networks)
scaler = StandardScaler()
X_train_resampled = scaler.fit_transform(X_train_resampled)
X_test = scaler.transform(X_test)

# Convert the data to PyTorch tensors
X_train_tensor = torch.tensor(X_train_resampled, dtype=torch.float32)
X_test_tensor = torch.tensor(X_test, dtype=torch.float32)
y_train_tensor = torch.tensor(y_train_resampled.values, dtype=torch.long)
y_test_tensor = torch.tensor(y_test.values, dtype=torch.long)

# Define the Deep Learning model
class ARFModel(nn.Module):
    def __init__(self, input_dim):
        super(ARFModel, self).__init__()
        self.layer11 = nn.Linear(input_dim, 128)
        self.batchnorm11 = nn.BatchNorm1d(128)
        self.layer1 = nn.Linear(128, 64)
        self.batchnorm1 = nn.BatchNorm1d(64)
        self.layer2 = nn.Linear(64, 32)
        self.batchnorm2 = nn.BatchNorm1d(32)
        self.layer3 = nn.Linear(32, 16)
        self.batchnorm3 = nn.BatchNorm1d(16) # Batch normalization
        self.output = nn.Linear(16, 2)
        self.dropout = nn.Dropout(0.3) # Dropout layer to reduce overfitting

    def forward(self, x):
        x = F.relu(self.batchnorm11(self.layer11(x)))
        x = self.dropout(x)
        x = F.relu(self.batchnorm1(self.layer1(x)))
        x = self.dropout(x)
        x = F.relu(self.batchnorm2(self.layer2(x)))
        x = self.dropout(x)
        x = F.relu(self.batchnorm3(self.layer3(x)))
        x = self.dropout(x)
        x = self.output(x)
        return x

# Initialize model, loss function, and optimizer
input_dim = X_train_tensor.shape[1]
model = ARFModel(input_dim=input_dim)

# Compute class weights to handle imbalance in the dataset
class_weights = compute_class_weight('balanced', classes=np.array([0, 1]), y=y_train_resampled)
class_weights = torch.tensor(class_weights, dtype=torch.float32)

# Define the loss function (CrossEntropyLoss) with class weights
criterion = nn.CrossEntropyLoss(weight=class_weights)
optimizer = optim.AdamW(model.parameters(), lr=0.001)
```

Step 6: Evaluate Neural Network Model(Continued)...

```
# Training loop (200 epochs)
num_epochs = 200
for epoch in range(num_epochs):
    model.train()
    optimizer.zero_grad()
    outputs = model(X_train_tensor)
    loss = criterion(outputs, y_train_tensor)
    loss.backward()
    optimizer.step()

    # Print the loss every 10 epochs
    if (epoch + 1) % 10 == 0:
        print(f"Epoch [{epoch+1}/{num_epochs}], Loss: {loss.item():.4f}")

# Evaluate the model on the test set
model.eval()
with torch.no_grad():
    outputs = model(X_test_tensor)
    _, predicted = torch.max(outputs, 1)

# Calculate various evaluation metrics
accuracy = accuracy_score(y_test_tensor, predicted)
y_prob = torch.softmax(outputs, dim=1)[: , 1]
roc_auc = roc_auc_score(y_test_tensor, y_prob)
precision = precision_score(y_test_tensor, predicted)
recall = recall_score(y_test_tensor, predicted)
f1 = f1_score(y_test_tensor, predicted)

# Print Model Performance Metrics
cf = classification_report(y_test_tensor, predicted)
cm = confusion_matrix(y_test_tensor, predicted)
print_model_performance_metrics('Neural Network', accuracy, auc_roc, precision, recall, f1, cf , cm)
```

Neural Network Performance:

Accuracy: 0.7889
AUC-ROC: 0.5423
Precision: 0.7271
Recall: 0.3461
F1-Score: 0.4689

Classification Report:

	precision	recall	f1-score	support
0	0.80	0.95	0.87	2485
1	0.73	0.35	0.47	916
accuracy			0.79	3401
macro avg	0.76	0.65	0.67	3401
weighted avg	0.78	0.79	0.76	3401

Confusion Matrix for Neural Network:

```
[[2366 119]
 [ 599 317]]
```

Here ARFModel is evaluated and following are observations,

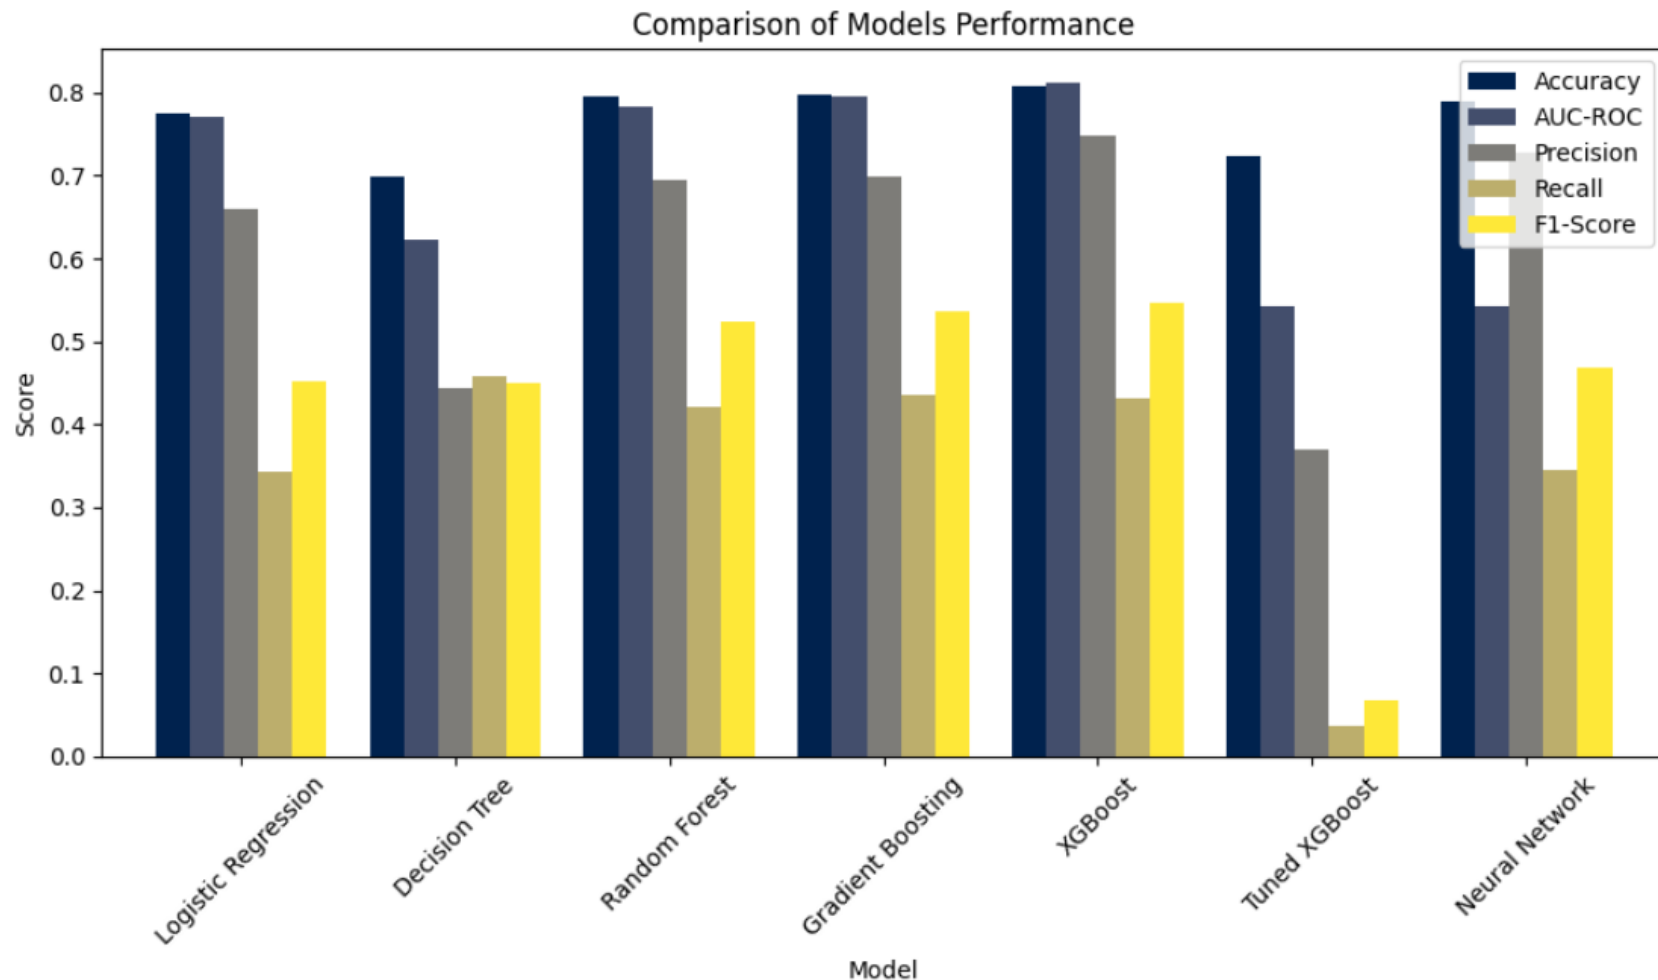
The Neural Network model achieved 78.8% accuracy, but its AUC-ROC (0.5423) and recall (32.5%) indicate weak discrimination in predicting mortality. It correctly classifies survival cases well (precision: 74.1%) but struggles to identify mortality cases.

Compared to XGBoost before tuning (accuracy: 80.8%, recall: ~43%), the Neural Network underperforms in recall and overall discrimination. However, unlike XGBoost after tuning, which saw a drop in accuracy (72.4%) and recall (3.7%), the Neural Network maintains a relatively stable performance.

Step 7: Visualize all models metrics

The following plot visually compares all models, showcasing their performance across key metrics such as accuracy, AUC-ROC, precision, recall, and F1-score.

```
#Show comparision of all the models  
plot_model_metrics_comparison(metrics_df)
```



Conclusion :

Among all the models evaluated, XGBoost demonstrated the best performance in predicting mortality in ARF patients. Despite hyperparameter tuning, its initial version outperformed other models, including the neural network. The feature analysis suggests that laboratory test results played a significant role in prediction, indicating their importance in assessing ARF severity and patient outcomes. Future improvements could involve further feature engineering, advanced ensemble methods, or incorporating temporal trends in lab values for better predictive accuracy.

Future improvements:

Future improvements could involve further feature engineering, additional hyperparameter tuning, or exploring different models to enhance predictive accuracy. Incorporating temporal trends in lab values and leveraging ensemble methods could also improve performance.